

eX3 workflows

Collecting best practices

moderated by Thomas Roehr

General constraints

- one project - multiple users
- heterogeneous target architectures (amd64 vs. arm64)
- shared resources in general
- popular and unpopular partitions / queues

Observation

- initial focus on ‘user’-centric ‘works-for-me’ approach
- initial permission setup (default umask) was not accounting for same user/group permissions
- docker - now disabled for majority of nodes due to security issues
- managing custom software stack via multiple / competing options:
 - module: <https://curc.readthedocs.io/en/latest/compute/modules.html>
 - conda: <https://docs.conda.io/en/latest/>
 - docker: <https://docs.docker.com/engine/reference/builder/>
 - spack: <https://spack.io/>
 - ...

Module

- `module avail <optional-pattern>`
 - gives you a list of all installed modules
- `module list`
 - show loaded modules
- `module show <name-of-module>`
 - gives you the list of module specific environment setup
- `module (un)load <name-of-module>`

Modules (Custom)

Writing custom module files, e.g., in `local_model`:

1. Compile your component
2. Write your own module file as 'custom-module'
3. Activate module
 - `module use <user-folder>`
 - append folder <user-folder> to load directories
 - `module load custom-module`

Observation

- filesystem usage
 - no dedicated backup of data - so every use might come with a different approach
 - syncing from Google Drive data bound to individual user accounts
 - no default folder structure, e.g. currently trying to seek a reasonable structure
- GPU support
 - nvidia-toolkit (will be preinstalled as module)
 - available, yet somewhat unclear what versions to use
 - starting point to monitor your GPU usage: nvidia-smi
 - <https://developer.nvidia.com/nvidia-system-management-interface>
- Running SLURM jobs with sbatch - good to know some details

SLURM - Workload Manager <https://slurm.schedmd.com/overview.html>

Using sbatch to send jobs (scripts) to the job queue:

- `sbatch -A <user> -o <job.log> -J <jobname> -p <partition> --mincpus=<number-of-cpus> --time=0-00:10:00 myjob.sh`

Options can also be encoded in the myjob.sh

```
#!/bin/bash
#SBATCH --job-name=myjob
#SBATCH --gres=gpu:1
```

For an interactive usage:

- `srun -p <partition> --mincpus 128 --time 0-00:10:00 --pty bash`

Manual allocation:

- `salloc -N1 xterm`
- depending on the usage of salloc you might have to cancel with: `scancel <job-id>`

SLURM - Workload Manager

Monitoring / Inspecting jobs:

- `squeue --user=<user>`
 - show all jobs of selected user
- `scontrol show job <job-id>`
 - list details of the particular job
- `scontrol show reservations`
 - Your job might conflict with current reservations so that you might have to adapt the `--time` option
 - when you own the reservation: `sbatch --reservation=<reservation-id>`
 - there are also [‘magnetic’ reservations](#) where you will not need this

Potentially better ...

- Keep user experience as close as possible to their single machine experience
 - start software as it is run on a single test machine
- Scripting (Bash/Python/...)
- Standard folder structure for projects
 - data
 - logs
 - software
 - scripts
 - venvs
 - ..
- Usage of README.md to point to relevant documentation

Links

- eX3: <http://wiki.ex3.simula.no/doku.php?id=start>
- SLURM:
 - <https://slurm.schedmd.com/overview.html>
 - <https://curc.readthedocs.io/en/latest/running-jobs/slurm-commands.html>
- Module: <https://curc.readthedocs.io/en/latest/compute/modules.html>
- Nvidia SMI: <https://developer.nvidia.com/nvidia-system-management-interface>